

✨ AI SDK 4.0 is now available!

[Read Announcement](#)

> Menu

AI SDK Core > streamText

streamText()

Streams text generations from a language model.

You can use the `streamText` function for interactive use cases such as chat bots and other real-time applications. You can also generate UI components with tools.

```
1 import { openai } from '@ai-sdk/openai';
2 import { streamText } from 'ai';
3
4 const { textStream } = streamText({
5   model: openai('gpt-4-turbo'),
6   prompt: 'Invent a new holiday and describe its traditions.',
7 });
8
9 for await (const textPart of textStream) {
10   process.stdout.write(textPart);
11 }
```

To see `streamText` in action, check out [these examples](#).

Import

```
import { streamText } from "ai"
```



API Signature

Parameters

model: `LanguageModel`

The language model to use. Example: `openai('gpt-4-turbo')`

system: `string`

The system prompt to use that specifies the behavior of the model.

prompt: `string`

The input prompt to generate the text from.

messages: `Array<CoreSystemMessage | CoreUserMessage | CoreAssistantMessage | CoreToolMessage> | Array<UIMessage>`

A list of messages that represent a conversation. Automatically converts UI messages from the `useChat` hook.

CoreSystemMessage

role: `'system'`

The role for the system message.

content: `string`

The content of the message.

CoreUserMessage

role: `'user'`

The role for the user message.

content: `string | Array<TextPart | ImagePart | FilePart>`

The content of the message.

TextPart

type: `'text'`

The type of the message part.

text: `string`

The text content of the message part.

ImagePart

type: `'image'`

The type of the message part.

image: string | Uint8Array | Buffer | ArrayBuffer | URL

The image content of the message part. String are either base64 encoded content, base64 data URLs, or http(s) URLs.

mimeType?: string

The mime type of the image. Optional.

FilePart

type: 'file'

The type of the message part.

data: string | Uint8Array | Buffer | ArrayBuffer | URL

The file content of the message part. String are either base64 encoded content, base64 data URLs, or http(s) URLs.

mimeType: string

The mime type of the file.

CoreAssistantMessage

role: 'assistant'

The role for the assistant message.

content: string | Array<TextPart | ToolCallPart>

The content of the message.

TextPart

type: 'text'

The type of the message part.

text: string

The text content of the message part.

ToolCallPart

type: 'tool-call'

The type of the message part.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

args: object based on zod schema

Parameters generated by the model to be used by the tool.

CoreToolMessage

role: 'tool'

The role for the assistant message.

content: Array<ToolResultPart>

The content of the message.

ToolResultPart

type: 'tool-result'

The type of the message part.

toolCallId: string

The id of the tool call the result corresponds to.

toolName: string

The name of the tool the result corresponds to.

result: unknown

The result returned by the tool after execution.

isError?: boolean

Whether the result is an error or an error message.

tools: Record<string, CoreTool>

Tools that are accessible to and can be called by the model. The model needs to support calling tools.

CoreTool

description?: string

Information about the purpose of the tool including details on how and when it can be used by the model.

parameters: Zod Schema | JSON Schema

The schema of the input that the tool expects. The language model will use this to generate the input. It is also used to validate the output of the language model. Use descriptions to make the input understandable for the language model. You can either pass in a Zod schema or a JSON schema (using the `jsonSchema` function).

execute?: async (parameters: T, options: ToolExecutionOptions) ⇒ RESULT

An async function that is called with the arguments from the tool call and produces a result. If not provided, the tool will not be executed automatically.

ToolExecutionOptions

toolCallId: string

The ID of the tool call. You can use it e.g. when sending tool-call related information with stream data.

messages: CoreMessage[]

Messages that were sent to the language model to initiate the response that contained the tool call. The messages do not include the system prompt nor the assistant response that contained the tool call.

abortSignal: AbortSignal

An optional abort signal that indicates that the overall operation should be aborted.

toolChoice?: "auto" | "none" | "required" | { "type": "tool", "toolName": string }

The tool choice setting. It specifies how tools are selected for execution. The default is "auto". "none" disables tool execution. "required" requires tools to be executed. { "type": "tool", "toolName": string } specifies a specific tool to execute.

maxTokens?: number

Maximum number of tokens to generate.

temperature?: number

Temperature setting. The value is passed through to the provider. The range depends on the provider and model. It is recommended to set either `temperature` or `topP`, but not both.

topP?: number

Nucleus sampling. The value is passed through to the provider. The range depends on the provider and model. It is recommended to set either `temperature` or `topP`, but not both.

topK?: number

Only sample from the top K options for each subsequent token. Used to remove "long tail" low probability responses. Recommended for advanced use cases only. You usually only need to use temperature.

presencePenalty?: number

Presence penalty setting. It affects the likelihood of the model to repeat information that is already in the prompt. The value is passed through to the provider. The range depends on the provider and model.

frequencyPenalty?: number

Frequency penalty setting. It affects the likelihood of the model to repeatedly use the same words or phrases. The value is passed through to the provider. The range depends on the provider and model.

stopSequences?: string[]

Sequences that will stop the generation of the text. If the model generates any of these sequences, it will stop generating further text.

seed?: number

The seed (integer) to use for random sampling. If set and supported by the model, calls will generate deterministic results.

maxRetries?: number

Maximum number of retries. Set to 0 to disable retries. Default: 2.

abortSignal?: AbortSignal

An optional abort signal that can be used to cancel the call.

headers?: Record<string, string>

Additional HTTP headers to be sent with the request. Only applicable for HTTP-based providers.

maxSteps?: number

Maximum number of sequential LLM calls (steps), e.g. when you use tool calls. A maximum number is required to prevent infinite loops in the case of misconfigured tools. By default, it is set to 1.

experimental_continueSteps?: boolean

Enable or disable continue steps. Disabled by default.

experimental_telemetry?: TelemetrySettings

Telemetry configuration. Experimental feature.

TelemetrySettings

isEnabled?: boolean

Enable or disable telemetry. Disabled by default while experimental.

recordInputs?: boolean

Enable or disable input recording. Enabled by default.

recordOutputs?: boolean

Enable or disable output recording. Enabled by default.

functionId?: string

Identifier for this function. Used to group telemetry data by function.

metadata?: Record<string, string | number | boolean | Array<null | undefined | string> | Array<null | undefined | number> | Array<null | undefined | boolean>>

Additional information to include in the telemetry data.

experimental_toolCallStreaming?: boolean

Enable streaming of tool call deltas as they are generated. Disabled by default.

`experimental_providerMetadata?: Record<string, Record<string, JSONValue>> | undefined`

Optional metadata from the provider. The outer key is the provider name. The inner values are the metadata. Details depend on the provider.

`experimental_activeTools?: Array<TOOLNAME> | undefined`

The tools that are currently active. All tools are active by default.

`experimental_repairToolCall?: (options: ToolCallRepairOptions) => Promise<LanguageModelV1FunctionToolCall | null>`

A function that attempts to repair a tool call that failed to parse. Return either a repaired tool call or null if the tool call cannot be repaired.

ToolCallRepairOptions

`system: string | undefined`
The system prompt.

`messages: CoreMessage[]`
The messages in the current generation step.

`toolCall: LanguageModelV1FunctionToolCall`
The tool call that failed to parse.

`tools: TOOLS`
The tools that are available.

`parameterSchema: (options: { toolName: string }) => JSONSchema7`
A function that returns the JSON Schema for a tool.

`error: NoSuchToolError | InvalidToolArgumentsError`
The error that occurred while parsing the tool call.

`onChunk?: (event: OnChunkResult) => Promise<void> | void`

Callback that is called for each chunk of the stream. The stream processing will pause until the callback promise is resolved.

OnChunkResult

`chunk: TextStreamPart`
The chunk of the stream.

TextStreamPart

`type: 'text-delta'`
The type to identify the object as text delta.

`textDelta: string`

The text delta.

TextStreamPart

type: 'tool-call'

The type to identify the object as tool call.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

args: object based on zod schema

Parameters generated by the model to be used by the tool.

TextStreamPart

type: 'tool-call-streaming-start'

Indicates the start of a tool call streaming. Only available when streaming tool calls.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

TextStreamPart

type: 'tool-call-delta'

The type to identify the object as tool call delta. Only available when streaming tool calls.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

argsTextDelta: string

The text delta of the tool call arguments.

TextStreamPart

type: 'tool-result'

The type to identify the object as tool result.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

args: object based on zod schema

Parameters generated by the model to be used by the tool.

result: any

The result returned by the tool after execution has completed.

onStepFinish?: (result: onStepFinishResult) => Promise<void> | void
Callback that is called when a step is finished.

onStepFinishResult

stepType: "initial" | "continue" | "tool-result"

The type of step. The first step is always an "initial" step, and subsequent steps are either "continue" steps or "tool-result" steps.

finishReason: "stop" | "length" | "content-filter" | "tool-calls" | "error" | "other" | "unknown"

The reason the model finished generating the text for the step.

usage: TokenUsage

The token usage of the step.

TokenUsage

promptTokens: number

The total number of tokens in the prompt.

completionTokens: number

The total number of tokens in the completion.

totalTokens: number

The total number of tokens generated.

text: string

The full text that has been generated.

toolCalls: ToolCall[]

The tool calls that have been executed.

toolResults: ToolResult[]

The tool results that have been generated.

warnings: Warning[] | undefined

Warnings from the model provider (e.g. unsupported settings).

response?: Response

Response metadata.

Response

id: string

The response identifier. The AI SDK uses the ID from the provider response when available, and generates an ID otherwise.

model: string

The model that was used to generate the response. The AI SDK uses the response model from the provider response when available, and the model from the function call otherwise.

timestamp: Date

The timestamp of the response. The AI SDK uses the response timestamp from the provider response when available, and creates a timestamp otherwise.

headers?: Record<string, string>

Optional response headers.

isContinued: boolean

True when there will be a continuation step with a continuation text.

experimental_providerMetadata?: Record<string, Record<string, JSONValue>> | undefined

Optional metadata from the provider. The outer key is the provider name. The inner values are the metadata. Details depend on the provider.

onFinish?: (result: OnFinishResult) \Rightarrow Promise<void> | void

Callback that is called when the LLM response and all request tool executions (for tools that have an `execute` function) are finished.

OnFinishResult

finishReason: "stop" | "length" | "content-filter" | "tool-calls" | "error" | "other" | "unknown"

The reason the model finished generating the text.

usage: TokenUsage

The token usage of the generated text.

TokenUsage

promptTokens: number

The total number of tokens in the prompt.

completionTokens: number

The total number of tokens in the completion.

totalTokens: number

The total number of tokens generated.

experimental_providerMetadata: Record<string, Record<string, JSONValue>> | undefined

Optional metadata from the provider. The outer key is the provider name. The inner values are the metadata. Details depend on the provider.

text: string

The full text that has been generated.

toolCalls: ToolCall[]

The tool calls that have been executed.

toolResults: ToolResult[]

The tool results that have been generated.

warnings: Warning[] | undefined

Warnings from the model provider (e.g. unsupported settings).

response?: Response

Response metadata.

Response

id: string

The response identifier. The AI SDK uses the ID from the provider response when available, and generates an ID otherwise.

model: string

The model that was used to generate the response. The AI SDK uses the response model from the provider response when available, and the model from the function call otherwise.

timestamp: Date

The timestamp of the response. The AI SDK uses the response timestamp from the provider response when available, and creates a timestamp otherwise.

headers?: Record<string, string>

Optional response headers.

messages: Array<CoreAssistantMessage | CoreToolMessage>

The response messages that were generated during the call. It consists of an assistant message, potentially containing tool calls. When there are tool results, there is an additional tool message with the tool results that are available. If there are tools that do not have execute functions, they are not included in the tool results and need to be added separately.

steps: Array<StepResult>

Response information for every step. You can use this to get information about intermediate steps, such as the tool calls or the response headers.

Returns

finishReason: Promise<'stop' | 'length' | 'content-filter' | 'tool-calls' | 'error' | 'other' | 'unknown'>

The reason why the generation finished. Resolved when the response is finished.

usage: Promise<CompletionTokenUsage>

The token usage of the generated text. Resolved when the response is finished.

CompletionTokenUsage

promptTokens: number

The total number of tokens in the prompt.

completionTokens: number

The total number of tokens in the completion.

totalTokens: number

The total number of tokens generated.

experimental_providerMetadata: Promise<Record<string, Record<string, JSONValue>> | undefined>

Optional metadata from the provider. Resolved when the response is finished. The outer key is the provider name. The inner values are the metadata. Details depend on the provider.

responseMessages: Promise<Array<CoreAssistantMessage | CoreToolMessage>>

The response messages that were generated during the call. It consists of an assistant message, potentially containing tool calls. When there are tool results, there is an additional tool message with the tool results that are available. If there are tools that do not have execute functions, they are not included in the tool results and need to be added separately. Resolved when the response is finished.

text: Promise<string>

The full text that has been generated. Resolved when the response is finished.

toolCalls: Promise<ToolCall[]>

The tool calls that have been executed. Resolved when the response is finished.

toolResults: Promise<ToolResult[]>

The tool results that have been generated. Resolved when the all tool executions are finished.

request?: Promise<RequestMetadata>

Request metadata.

RequestMetadata

body: string

Raw request HTTP body that was sent to the provider API as a string (JSON should be stringified).

response?: Promise<ResponseMetadata>

Response metadata. Resolved when the response is finished.

ResponseMetadata

id: string

The response identifier. The AI SDK uses the ID from the provider response when available, and generates an ID otherwise.

model: string

The model that was used to generate the response. The AI SDK uses the response model from the provider response when available, and the model from the function call otherwise.

timestamp: Date

The timestamp of the response. The AI SDK uses the response timestamp from the provider response when available, and creates a timestamp otherwise.

headers?: Record<string, string>

Optional response headers.

messages: Array<CoreAssistantMessage | CoreToolMessage>

The response messages that were generated during the call. It consists of an assistant message, potentially containing tool calls. When there are tool results, there is an additional tool message with the tool results that are available. If there are tools that do not have execute functions, they are not included in the tool results and need to be added separately.

warnings: Warning[] | undefined

Warnings from the model provider (e.g. unsupported settings).

steps: Promise<Array<StepResult>>

Response information for every step. You can use this to get information about intermediate steps, such as the tool calls or the response headers.

StepResult

stepType: "initial" | "continue" | "tool-result"

The type of step. The first step is always an "initial" step, and subsequent steps are either "continue" steps or "tool-result" steps.

text: string

The generated text by the model.

toolCalls: array

A list of tool calls made by the model.

toolResults: array

A list of tool results returned as responses to earlier tool calls.

finishReason: 'stop' | 'length' | 'content-filter' | 'tool-calls' | 'error' | 'other' | 'unknown'

The reason the model finished generating the text.

usage: CompletionTokenUsage

The token usage of the generated text.

CompletionTokenUsage

promptTokens: number

The total number of tokens in the prompt.

completionTokens: number

The total number of tokens in the completion.

totalTokens: number

The total number of tokens generated.

request?: RequestMetadata

Request metadata.

RequestMetadata

body: string

Raw request HTTP body that was sent to the provider API as a string (JSON should be stringified).

response?: ResponseMetadata

Response metadata.

ResponseMetadata

id: string

The response identifier. The AI SDK uses the ID from the provider response when available, and generates an ID otherwise.

model: string

The model that was used to generate the response. The AI SDK uses the response model from the provider response when available, and the model from the function call otherwise.

timestamp: Date

The timestamp of the response. The AI SDK uses the response timestamp from the provider response when available, and creates a timestamp otherwise.

headers?: Record<string, string>

Optional response headers.

messages: Array<CoreAssistantMessage | CoreToolMessage>

The response messages that were generated during the call. It consists of an assistant message, potentially containing tool calls. When there are tool results, there is an additional tool message with the tool results that are available. If there are tools that do not have execute functions, they are not included in the tool results and need to be added separately.

warnings: Warning[] | undefined

Warnings from the model provider (e.g. unsupported settings).

isContinued: boolean

True when there will be a continuation step with a continuation text.

experimental_providerMetadata?: Record<string, Record<string, JSONValue>> | undefined

Optional metadata from the provider. The outer key is the provider name. The inner values are the metadata. Details depend on the provider.

textStream: AsyncIterable<string> & ReadableStream<string>

A text stream that returns only the generated text deltas. You can use it as either an AsyncIterable or a ReadableStream. When an error occurs, the stream will throw the error.

fullStream: AsyncIterable<TextStreamPart> & ReadableStream<TextStreamPart>

A stream with all events, including text deltas, tool calls, tool results, and errors. You can use it as either an AsyncIterable or a ReadableStream. Only errors that stop the stream, such as network errors, are thrown.

TextStreamPart

type: 'text-delta'

The type to identify the object as text delta.

textDelta: string

The text delta.

TextStreamPart

type: 'tool-call'

The type to identify the object as tool call.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

args: object based on zod schema

Parameters generated by the model to be used by the tool.

TextStreamPart

type: 'tool-call-streaming-start'

Indicates the start of a tool call streaming. Only available when streaming tool calls.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

TextStreamPart

type: 'tool-call-delta'

The type to identify the object as tool call delta. Only available when streaming tool calls.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

argsTextDelta: string

The text delta of the tool call arguments.

TextStreamPart

type: 'tool-result'

The type to identify the object as tool result.

toolCallId: string

The id of the tool call.

toolName: string

The name of the tool, which typically would be the name of the function.

args: object based on zod schema

Parameters generated by the model to be used by the tool.

result: any

The result returned by the tool after execution has completed.

TextStreamPart

type: 'error'

The type to identify the object as error.

error: Error

Describes the error that may have occurred during execution.

TextStreamPart

type: 'step-finish'

The type to identify the object as step finish.

finishReason: 'stop' | 'length' | 'content-filter' | 'tool-calls' |
'error' | 'other' | 'unknown'

The reason the model finished generating the text.

usage: TokenUsage

The token usage of the generated text.

TokenUsage

promptTokens: number

The total number of tokens in the prompt.

completionTokens: number

The total number of tokens in the completion.

totalTokens: number

The total number of tokens generated.

isContinued: boolean

True when there will be a continuation step with a continuation text.

response?: Response

Response metadata.

Response

id: string

The response identifier. The AI SDK uses the ID from the provider response when available, and generates an ID otherwise.

model: string

The model that was used to generate the response. The AI SDK uses the response model from the provider response when available, and the model from the function call otherwise.

timestamp: Date

The timestamp of the response. The AI SDK uses the response timestamp from the provider response when available, and creates a timestamp otherwise.

TextStreamPart

type: 'finish'

The type to identify the object as finish.

finishReason: 'stop' | 'length' | 'content-filter' | 'tool-calls' | 'error' | 'other' | 'unknown'

The reason the model finished generating the text.

usage: TokenUsage

The token usage of the generated text.

TokenUsage

promptTokens: number

The total number of tokens in the prompt.

completionTokens: number

The total number of tokens in the completion.

totalTokens: number

The total number of tokens generated.

response?: Response

Response metadata.

Response

id: string

The response identifier. The AI SDK uses the ID from the provider response when available, and generates an ID otherwise.

model: string

The model that was used to generate the response. The AI SDK uses the response model from the provider response when available, and the model from the function call otherwise.

timestamp: Date

The timestamp of the response. The AI SDK uses the response timestamp from the provider response when available, and creates a timestamp otherwise.

```
pipeDataStreamToResponse: (response: ServerResponse, options:
    PipeDataStreamToResponseOptions } => void
```

Writes stream data output to a Node.js response-like object. It sets a `Content-Type` header to `text/plain; charset=utf-8` and writes each stream data part as a separate chunk.

PipeDataStreamToResponseOptions

status?: number

The response status code.

statusText?: string

The response status text.

headers?: HeadersInit

The response headers.

data?: StreamData

The stream data object.

getErrorMessage?: (error: unknown) => string

A function to get the error message from the error object. By default, all errors are masked as "" for safety reasons.

sendUsage?: boolean

Whether to send the usage information in the stream. Defaults to true.

```
pipeTextStreamToResponse: (response: ServerResponse, init?: ResponseInit =>
    void
```

Writes text delta output to a Node.js response-like object. It sets a `Content-Type` header to `text/plain; charset=utf-8` and writes each text delta as a separate chunk.

ResponseInit

status?: number

The response status code.

statusText?: string

The response status text.

headers?: Record<string, string>

The response headers.

toDataStream: (options?: ToDataStreamOptions) ⇒ Response

Converts the result to a data stream.

ToDataStreamOptions

data?: StreamData

The stream data object.

getErrorMessage?: (error: unknown) ⇒ string

A function to get the error message from the error object. By default, all errors are masked as "" for safety reasons.

sendUsage?: boolean

Whether to send the usage information in the stream. Defaults to true.

toDataStreamResponse: (options?: ToDataStreamResponseOptions) ⇒ Response

Converts the result to a streamed response object with a stream data part stream. It can be used with the `useChat` and `useCompletion` hooks.

ToDataStreamResponseOptions

status?: number

The response status code.

statusText?: string

The response status text.

headers?: Record<string, string>

The response headers.

data?: StreamData

The stream data object.

getErrorMessage?: (error: unknown) ⇒ string

A function to get the error message from the error object. By default, all errors are masked as "" for safety reasons.

sendUsage?: boolean

Whether to send the usage information in the stream. Defaults to true.

toTextStreamResponse: (init?: ResponseInit) ⇒ Response

Creates a simple text stream response. Each text delta is encoded as UTF-8 and sent as a separate chunk. Non-text-delta events are ignored.

ResponseInit

status?: number

The response status code.

statusText?: string


The response status text.


headers?: Record<string, string>


The response headers.

Examples

 [Learn to stream text generated by a language model in Next.js](#) ↗

 [Learn to stream chat completions generated by a language model in Next.js](#) ↗

 [Learn to stream text generated by a language model in Node.js](#) ↗

 [Learn to stream chat completions generated by a language model in Node.js](#) ↗

Previous

< generateText

Next

generateObject >



Resources

[Docs](#)

[Cookbook](#)

[Providers](#)

[Showcase](#)

[GitHub](#)

[Discussions](#)

More

[Playground](#)

[VX](#)

[Contact Sales](#)

About Vercel

[Next.js + Vercel](#)

[Open Source Software](#)

[GitHub](#)

[X](#)

Legal

[Privacy Policy](#)

© 2024 Vercel, Inc.

